

## Assignment 3 (Part II of II) Total Points 50

### Objectives

In this assignment you will extend a simple *command line interpreter* (or *shell*) developed by your instructor. This assignment will allow you to gain experience in the following areas:

- **Process Communication:** Extend your knowledge of processes (i.e. `fork`, `exec`, and `wait`) by including process communication functionality into your shell program.
- **More 'C' Programming:** This includes variable declaration, data types, arrays, pointers, operators, expressions, selection statements, looping statements, functions, structs, and header files.

### System and Standard Lib Functions

This assignment will use the system and standard library functions listed below. Please become familiar with the syntax and usage of these calls. Detailed information about each function listed below can be found using the `man` command from the console: i.e. `man pipe`, will show the man page (short for manual page) for the `pipe` function.

- **Process Communication:** `int pipe(int filedes[2])`
- **Duplicate File Descriptor:** `int dup(int fildes)`
- **Blocking Operation:** `pid_t waitpid(pid_t wpid, int *status, int options)`

### Provided Files

Modify the provided `hw2.c` and `shell.c` files. You cannot modify `shell.h`.

### Todo

To include `pipe()` and `dup()` functionality into your shell. Only the `int execute(command_t* p_cmd)` function in the `shell.c` file needs be modified. **Multiple pipe operations may be used in this assignment.** The `|` shell symbol represents a UNIX pipe operation. For instance, `ls | wc -l` is an example shell command that “pipes” the `stdout` of the `ls` shell command to the `wc` word count shell command. For further help, see the Additional Guidance section at the end of this document. Note, neither `hw2.c` nor `shell.h` need to be modified in this assignment.

### Collaboration and Plagiarism

This is an **individual assignment**, i.e. **no collaboration is permitted**. Plagiarism will not be tolerated. Submitted solutions that are very similar (determined by the instructor) will be given a grade of zero. **Please do your own work.**

## Submission

Create a compressed tarball, i.e. *tar.gz*, that only contains the completed *hw2.c*, *shell.h* and *shell.c* files. The name of the compressed tarball must be your last name in lower case. For example, *ritchie.tar.gz* would be correct if the original co-developer of UNIX (Dennis Ritchie) submitted the assignment. Only assignments submitted in the correct format will be accepted (no exceptions). Submit the compressed tarball (via OAKS) to the Dropbox setup for this assignment. You may resubmit the compressed tarball as many times as you like, only the most recent will be graded.

To be fair to everyone, late assignments will not be accepted. Exceptions will only be made for extenuating circumstances, i.e. death in the family, health related problems, etc. You will be given a week to complete this assignment. Poor time management is not excuse. Please do not email assignment after the due date, it will not be accepted. Please feel free to setup an appointment to discuss the assigned coding problem. I will be more than happy to listen to your approach and make suggestions. However, I cannot tell you how to code the solution. Additionally, code debugging is your job. You may use the debugger (gdb) or print statements to help understand why your solution is not working correctly, your choice.

## Grading Rubric

For this assignment the grading rubric is provided in the table shown below.

Program Compiles	10 points
Program Runs with no errors	10 points
execute() function implementation	30 points

In particular, the assignment will be graded as follows, if the submitted solution

- does not compile: 0 of 50 points
- compiles but does not run: 10 of 50 points
- compiles and runs with errors: 15 of 50 points
- compiles and runs without errors: 20 of 50 points
- all functions correctly implemented: 50 of 50 points

## Additional Guidance

### UNIX Pipe

Pipes are the oldest form of UNIX Inter-process communication (IPC) and are provided by all UNIX systems. They have two limitations:

1. They are half-duplex, i.e. data only flows in one direction.
2. They can be used only between processes that have a common ancestor. Normally a pipe is created by a process, that process calls a fork, and the pipe is used between the parent and child.

A pipe is created by calling the `int pipe(int fd[2])` function. Two file descriptors are returned through the `fd` argument: `fd[0]` is open for reading, and `fd[1]` is open for writing. Normally, the process that calls `pipe` then calls `fork`, creating an IPC channel from the parent to the child (or vice versa). Figure 1 illustrates this concept.

What happens after the `fork` depends on which direction of data flow we want. For a pipe from parent to the child, the parent closes the read end of the pipe (`fd[0]`) and the child closes the write end (`fd[1]`). For a pipe from the child to the parent, the parent closes `fd[1]` and the child closes `fd[0]`.




Figure 1: UNIX pipe concept between parent and child process.

## Duplicate File Descriptor

An existing file descriptor is duplicated using `int dup(int fildes)`. In general, if `fildes` is 0 then the standard in (`stdin`) is duplicated, and if `fildes` is 1 then the standard out is duplicated. These duplicated descriptors can be used interchangeably by the parent and child process.

## Example code that uses pipe and dup

An example code segment that illustrates how to code the equivalent `ls | wc -l` shell command. Please note, this example is not considered complete, and is only meant to guide you in this coding assignment.

```
int main(int argc, char** argv) {
    int child_process_status;
    int fds[2];
    pid_t cpid1, cpid2;
    char* pargs[] = { "ls", NULL };
    char* cargs[] = { "wc", "-l", NULL };

    pipe(fds);

    if ((cpid1 = fork()) == 0) {
        close(1);                /* close normal stdout */
        dup(fds[1]);              /* make stdout same as fds[1] */
        close(fds[0]);           /* we don't need this */
        execv("/bin/ls", pargs);
    }

    if ((cpid2 = fork()) == 0) {
        close(0);                /* close normal stdin */
        dup(fds[0]);              /* make stdin same as fds[0] */
        close(fds[1]);           /* we don't need this */
        execv("/usr/bin/wc", cargs);
    }

    close(fds[0]);
    close(fds[1]);
    waitpid(cpid1, &child_process_status, 0);
    waitpid(cpid2, &child_process_status, 0);
    return 0;
}
```